

## **Kubernetes Service Deployment using Helm Chart**

**Helm Chart** : A Helm chart is a package that contains all the resources needed to deploy an application to a Kubernetes cluster. Basically we can define as a package manager for kubernetes.

### **Steps to Create the Helm Chart**

**Create the Helm Chart:** To create a Helm chart, run the following command:

```
helm create <repo-name>
```

Example:

```
helm create nginx
```

**Navigate to the Chart Directory:** Once the chart is created, navigate to the chart directory:

```
cd nginx
```

1. Open this directory in your preferred text editor to modify the templates.
2. **Modify the Templates:** Modify the YAML template files inside the Helm chart to define the Kubernetes resources required for the Nginx deployment.

## Configuration Breakdown

### 1. `values.yaml`

The `values.yaml` file defines the configurable parameters for the Helm chart. Key settings include:

#### Replica Count and Deployment Name:

```
replicaCount: 1
deploymentName: test-nginx-helm
namespace: nginx
```

**Nginx Image Configuration:** The image repository is set to `nginx`, and the `IfNotPresent` pull policy ensures that the image is pulled only if it isn't already available on the node.

```
image:
  repository: nginx
  tag: latest
  pullPolicy: IfNotPresent
```

**Service Configuration:** The service is of type `ClusterIP` and exposes port 80.

```
service:
  type: ClusterIP
  port: 80
```

**Ingress Configuration:** The ingress is configured with AWS NLB annotations, enabling secure HTTPS traffic using TLS with a self-signed certificate.

#### ingress:

```
ingressName: nginx-nlb-ingress
enabled: true
className: nlb
annotations:
  kubernetes.io/ingress.class: "nlb"
  service.beta.kubernetes.io/aws-load-balancer-type: "nlb"
  nginx.ingress.kubernetes.io/rewrite-target: /
  nginx.ingress.kubernetes.io/ssl-redirect: 'true'
  service.beta.kubernetes.io/aws-load-balancer-proxy-protocol: "*"
  service.beta.kubernetes.io/aws-load-balancer-nlb-target-type: "ip"
  service.beta.kubernetes.io/aws-load-balancer-scheme: internet-facing
  service.beta.kubernetes.io/aws-load-balancer-backend-protocol: tcp
  acme.cert-manager.io/http01-edit-in-place: 'true'
  cert-manager.io/cluster-issuer: letsencrypt-prod
  kubernetes.io/tls-acme: 'true'
hosts:
  - host: test-nginx.example.com
    paths:
      - path: /
        pathType: prefix
tls:
  - secretName: chart-example-tls
    hosts:
      - test-nginx.example.com
```

**Horizontal Pod Autoscaling (HPA):** HPA is enabled to adjust the number of replicas based on CPU utilization (with a target of 70%).

```
autoscaling:
  enabled: true
  minReplicas: 1
  maxReplicas: 3
  targetCPUUtilizationPercentage: 70
```

**Resource Allocation:** Requests and limits for CPU and memory resources are defined.

```
resources:
  requests:
    memory: "128Mi"
    cpu: "100m"
  limits:
    memory: "128Mi"
    cpu: "100m"
```

## Kubernetes Resource Templates

The Helm chart defines several Kubernetes resources, which are dynamically populated from the `values.yaml` file:

Namespace: `namespace.yaml`

This file defines the Kubernetes namespace for the deployment:

```
apiVersion: v1
kind: Namespace
metadata:
  name: "{{ .Values.namespace }}"
```

Deployment: **deployment.yaml**

This file defines the Nginx deployment, which is tied to the values specified in **values.yaml**. If autoscaling is enabled, the replica count will adjust dynamically. Otherwise, it will use the static **replicaCount**.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ .Values.deploymentName }}
  namespace: {{ .Values.namespace }}
  labels:
    {{- include "nginx.labels" . | nindent 4 }}
spec:
  {{- if not .Values.autoscaling.enabled }}
  replicas: {{ .Values.replicaCount }}
  {{- end }}
  selector:
    matchLabels:
      {{- include "nginx.selectorLabels" . | nindent 6 }}
  template:
    metadata:
      {{- with .Values.podAnnotations }}
      annotations:
        {{- toYaml . | nindent 8 }}
      {{- end }}
      labels:
        {{- include "nginx.labels" . | nindent 8 }}
        {{- with .Values.podLabels }}
        {{- toYaml . | nindent 8 }}
```

```

        {{- end }}
spec:
  containers:
    - name: {{ .Chart.Name }}
      image: "{{ .Values.image.repository }}:{{ .Values.image.tag |
default .Chart.AppVersion }}"
      ports:
        - name: http
          containerPort: {{ .Values.service.port }}
          protocol: TCP
      resources:
        {{- toYaml .Values.resources | nindent 12 }}

```

### Service: service.yaml

This file defines a **ClusterIP** service that exposes Nginx to the cluster on port 80.

```

apiVersion: v1
kind: Service
metadata:
  name: {{ .Values.depServiceName }}
  namespace: {{ .Values.namespace }}
  labels:
    {{- include "nginx.labels" . | nindent 4 }}
spec:
  type: {{ .Values.service.type }}
  ports:
    - port: {{ .Values.service.port }}
      targetPort: http
      protocol: TCP
      name: http
  selector:
    {{- include "nginx.selectorLabels" . | nindent 4 }}

```

## Horizontal Pod Autoscaler (HPA): `hpa.yaml`

If autoscaling is enabled, the HPA adjusts the number of replicas based on CPU utilization.

```
{{- if .Values.autoscaling.enabled }}
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: {{ .Values.hpaName }}
  namespace: {{ .Values.namespace }}
  labels:
    {{- include "nginx.labels" . | nindent 4 }}
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: {{ .Values.deploymentName }}
  minReplicas: {{ .Values.autoscaling.minReplicas }}
  maxReplicas: {{ .Values.autoscaling.maxReplicas }}
  metrics:
    {{- if .Values.autoscaling.targetCPUUtilizationPercentage }}
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: {{
.Values.autoscaling.targetCPUUtilizationPercentage }}
    {{- end }}
{{- end }}
```

## Ingress: ingress.yaml

This file configures the Nginx ingress with AWS NLB annotations and enables TLS using a secret for secure HTTPS traffic.

```
{{- if .Values.ingress.enabled -}}
{{- $fullName := include "nginx.fullname" . -}}
{{- $svcPort := .Values.service.port -}}
{{- if and .Values.ingress.className (not (semverCompare ">=1.18-0"
.Capabilities.KubeVersion.GitVersion)) }}
  {{- if not (hasKey .Values.ingress.annotations
"kubernetes.io/ingress.class") }}
    {{- $_ := set .Values.ingress.annotations "kubernetes.io/ingress.class"
.Values.ingress.className}}
  {{- end }}
{{- end }}
{{- if semverCompare ">=1.19-0" .Capabilities.KubeVersion.GitVersion -}}
apiVersion: networking.k8s.io/v1
{{- else if semverCompare ">=1.14-0" .Capabilities.KubeVersion.GitVersion
-}}
apiVersion: networking.k8s.io/v1beta1
{{- else -}}
apiVersion: extensions/v1beta1
{{- end }}
kind: Ingress
metadata:
  name: {{ .Values.ingress.ingressName }}
  namespace: {{ .Values.namespace }}
  labels:
    {{- include "nginx.labels" . | nindent 4 }}
    {{- with .Values.ingress.annotations }}
  annotations:
    {{- toYaml . | nindent 4 }}
    {{- end }}
spec:
  {{- if and .Values.ingress.className (semverCompare ">=1.18-0"
.Capabilities.KubeVersion.GitVersion) }}
  ingressClassName: {{ .Values.ingress.className }}
  {{- end }}
```



```
{{- if .Values.ingress.tls }}
tls:
  {{- range .Values.ingress.tls }}
  - hosts:
    {{- range .hosts }}
    - {{ . | quote }}
    {{- end }}
    secretName: {{ .secretName }}
  {{- end }}
{{- end }}
rules:
  {{- range .Values.ingress.hosts }}
  - host: {{ .host | quote }}
    http:
      paths:
        {{- range .paths }}
        - path: {{ .path }}
          {{- if and .pathType (semverCompare ">=1.18-0"
$.Capabilities.KubeVersion.GitVersion) }}
          pathType: {{ .pathType }}
          {{- end }}
          backend:
            {{- if semverCompare ">=1.19-0"
$.Capabilities.KubeVersion.GitVersion }}
            service:
              name: {{ .Values.depServiceName }}
              port:
                number: {{ .Values.service.port }}
            {{- else }}
            serviceName: {{ $fullName }}
            servicePort: {{ $svcPort }}
            {{- end }}
          {{- end }}
        {{- end }}
      {{- end }}
    {{- end }}
```

---

## AWS Load Balancer Controller Integration

To integrate the AWS Load Balancer with the Nginx ingress, the **AWS Load Balancer Controller** should be deployed in the Kubernetes cluster. This controller automatically provisions and manages an NLB, ensuring that traffic is routed to the Nginx service according to the ingress rules.

### Annotations for NLB Integration:

- **service.beta.kubernetes.io/aws-load-balancer-type:**  
**"nlb"**: Specifies that the load balancer type should be an NLB.
- **service.beta.kubernetes.io/aws-load-balancer-scheme:**  
**internet-facing**: Configures the NLB to be internet-facing.
- **service.beta.kubernetes.io/aws-load-balancer-backend-protocol:** **tcp**: Defines the backend protocol for the NLB as TCP.
- **service.beta.kubernetes.io/aws-load-balancer-proxy-protocol:** **"\*"**: Enables proxy protocol support for the NLB.

To install the AWS Load Balancer Controller via Helm, use the following command:

```
helm install aws-load-balancer-controller  
eks/aws-load-balancer-controller
```

This controller will automatically create and manage the NLB and ensure that traffic is routed to the Nginx service as defined in the ingress configuration.

## To install helm chart from local helm directory

```
helm install nginx-test ./nginx
```

### Explanation:

- **helm install**: This command is used to install a Helm chart.
- **nginx-test**: This is the name you want to assign to the release. It is the name that will be used to reference the deployed chart.
- **./nginx**: This specifies the path to the chart. In this case, nginx is the chart directory located in the current directory (denoted by ./).

### Breakdown:

- **nginx-test**: The name of the Helm release.
- **./nginx**: Refers to the nginx chart, located in the current directory.

This command installs the Helm chart located in the nginx/ directory and names the release nginx-test. The ./ indicates that the chart is in the current working directory.

## To check the chart is install or not

```
helm list -n <namespace>
```